

```
import warnings
warnings.filterwarnings('ignore')

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.metrics import confusion_matrix, roc_auc_score, log_loss
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.ensemble import GradientBoostingClassifier

pd.set_option('display.max_columns', None)

df = pd.read_csv("/Bank Data.csv")
df.head()
```

	ID	Customer_ID	Month	Name	Age	SSN	Occupation	Annual_Income	Mo
0	0x160a	CUS_0xd40	September	Aaron Maashoh	23	821-00-0265	Scientist	19114.12	
1	0x160b	CUS_0xd40	October	Aaron Maashoh	24	821-00-0265	Scientist	19114.12	
2	0x160c	CUS_0xd40	November	Aaron Maashoh	24	821-00-0265	Scientist	19114.12	
3	0x160d	CUS_0xd40	December	Aaron Maashoh	24	821-00-0265	Scientist	19114.12	
4	0x1616	CUS_0x21b1	September	Rick Rothackerj	28	004-07-5839	_____	34847.84	

```
df.isnull().sum()

ID                0
Customer_ID       0
Month             0
Name             5015
Age              0
SSN              0
Occupation        0
Annual_Income     0
Monthly_Inhand_Salary  7498
Num_Bank_Accounts  0
Num_Credit_Card   0
Interest_Rate     0
Num_of_Loan       0
Type_of_Loan      5704
Delay_from_due_date  0
Num_of_Delayed_Payment  3498
Changed_Credit_Limit  0
Num_Credit_Inquiries  1035
Credit_Mix        0
Outstanding_Debt   0
Credit_Utilization_Ratio  0
Credit_History_Age  4470
Payment_of_Min_Amount  0
Total_EMI_per_month  0
Amount_invested_monthly  2271
Payment_Behaviour  0
Monthly_Balance    562
dtype: int64
```

```
df.drop(columns=['ID', 'Customer_ID', 'SSN', 'Name', 'Type_of_Loan', 'Payment_Behaviour'], inplace=True)
df.head()
```

	Month	Age	Occupation	Annual_Income	Monthly_Inhand_Salary	Num_Bank_Account
0	September	23	Scientist	19114.12	1824.843333	
1	October	24	Scientist	19114.12	1824.843333	
2	November	24	Scientist	19114.12	1824.843333	
3	December	24_	Scientist	19114.12	NaN	
4	September	28	_____	34847.84	3037.986667	

```
df.shape
```

(50000, 21)

```
df.Credit_Mix.value_counts().head(10)
```

Standard 18379
Good 12260
_ 9805
Bad 9556
Name: Credit_Mix, dtype: int64

```
# Assuming df is your DataFrame
df['Credit_Mix'].replace('_', 'neutral', inplace=True)
df['Credit_Mix'].fillna('neutral', inplace=True)
```

```
df.Credit_Mix.value_counts().head(10)
```

Standard 18379
Good 12260
neutral 9805
Bad 9556
Name: Credit_Mix, dtype: int64

```
df.isnull().sum()
```

Month 0
Age 0
Occupation 0
Annual_Income 0
Monthly_Inhand_Salary 7498
Num_Bank_Accounts 0
Num_Credit_Card 0
Interest_Rate 0
Num_of_Loan 0
Delay_from_due_date 0
Num_of_Delayed_Payment 3498
Changed_Credit_Limit 0
Num_Credit_Inquiries 1035
Credit_Mix 0
Outstanding_Debt 0
Credit_Utilization_Ratio 0
Credit_History_Age 4470
Payment_of_Min_Amount 0
Total_EMI_per_month 0
Amount_invested_monthly 2271
Monthly_Balance 562
dtype: int64

```
df.Delay_from_due_date.value_counts().sample(40)
```

55 296
39 250
25 1289
11 1573
2 669
9 1656
20 1300
-3 59
37 267
15 1759
62 279
31 397

```

53      291
41      307
48      363
17     1183
50      270
38      306
56      283
18     1335
34      320
-5        18
14     1636
52      302
10     1645
58      282
67         7
12     1625
59      250
29     1223
13     1761
57      269
3       848
66        12
1       668
45      269
4       825
22     1135
-2        71
54      308
Name: Delay_from_due_date, dtype: int64

```

Convert String to Float of Credit_History_Age column also fill the null value with it's mean

After doing a lot of analysis I have found all 'Credit_History_Age' null cause I have to convert all of str to float/int value. For this operation 'Credit_History_Age' columns became null. But I need this column. So that I'm do it preprocessing. I have taken only integer value from this columns that would be standard

```

def extract_age(row):
    if isinstance(row, str):
        years, months = 0, 0
        if 'Years' in row:
            years = int(row.split(' Years')[0])
        if 'Months' in row:
            months = int(row.split('and ')[1].split(' Months')[0])
        return years + months / 12
    else:
        return np.nan

df['Credit_History_Age'] = df['Credit_History_Age'].apply(extract_age)

# Fill missing values with the mean
mean_age = df['Credit_History_Age'].mean()
df['Credit_History_Age'].fillna(mean_age, inplace=True)

df.head(5)

```

	Month	Age	Occupation	Annual_Income	Monthly_Inhand_Salary	Num_Bank_Account
0	September	23	Scientist	19114.12	1824.843333	
1	October	24	Scientist	19114.12	1824.843333	
2	November	24	Scientist	19114.12	1824.843333	
3	December	24_	Scientist	19114.12	NaN	
4	September	28	_____	34847.84	3037.986667	

```

df.Payment_of_Min_Amount.value_counts()

Yes      26158
No       17849
NM        5993
Name: Payment_of_Min_Amount, dtype: int64

```

```

df.Occupation.value_counts()

_____      3438
Lawyer       3324
Engineer     3212
Architect    3195

```

```

Mechanic      3168
Developer     3146
Accountant    3133
Media_Manager 3130
Scientist     3104
Teacher       3103
Entrepreneur  3103
Journalist    3037
Doctor        3027
Manager       3000
Musician      2947
Writer        2933
Name: Occupation, dtype: int64

```

```
import numpy as np
```

```

# Replace "_____" with NaN
df['Occupation'].replace("_____", np.nan, inplace=True)

```

```

# Get the list of occupations
occupations_list = ["Lawyer", "Engineer", "Architect", "Mechanic", "Developer", "Accountant",
                    "Media_Manager", "Scientist", "Teacher", "Entrepreneur", "Journalist",
                    "Doctor", "Manager", "Musician", "Writer"]

```

```

# Replace NaN values with random values from the list
nan_mask = df['Occupation'].isnull()
df.loc[nan_mask, 'Occupation'] = np.random.choice(occupations_list, size=nan_mask.sum())

```

```
df.Occupation.value_counts()
```

```

Lawyer      3556
Architect   3436
Mechanic    3423
Engineer    3411
Accountant  3371
Developer   3365
Media_Manager 3356
Scientist   3352
Entrepreneur 3334
Teacher     3322
Journalist  3250
Doctor      3244
Manager     3234
Musician    3188
Writer      3158
Name: Occupation, dtype: int64

```

```
df.isnull().sum()
```

```

Month      0
Age         0
Occupation  0
Annual_Income  0
Monthly_Inhand_Salary  7498
Num_Bank_Accounts  0
Num_Credit_Card  0
Interest_Rate  0
Num_of_Loan  0
Delay_from_due_date  0
Num_of_Delayed_Payment  3498
Changed_Credit_Limit  0
Num_Credit_Inquiries  1035
Credit_Mix  0
Outstanding_Debt  0
Credit_Utilization_Ratio  0
Credit_History_Age  0
Payment_of_Min_Amount  0
Total_EMI_per_month  0
Amount_invested_monthly  2271
Monthly_Balance  562
dtype: int64

```

```
df.shape
```

```
(50000, 21)
```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype

```

```

---
0  Month                    50000 non-null object
1  Age                     50000 non-null object
2  Occupation              50000 non-null object
3  Annual_Income           50000 non-null object
4  Monthly_Inhand_Salary   42502 non-null float64
5  Num_Bank_Accounts       50000 non-null int64
6  Num_Credit_Card         50000 non-null int64
7  Interest_Rate           50000 non-null int64
8  Num_of_Loan             50000 non-null object
9  Delay_from_due_date     50000 non-null int64
10 Num_of_Delayed_Payment  46502 non-null object
11 Changed_Credit_Limit    50000 non-null object
12 Num_Credit_Inquiries    48965 non-null float64
13 Credit_Mix              50000 non-null object
14 Outstanding_Debt        50000 non-null object
15 Credit_Utilization_Ratio 50000 non-null float64
16 Credit_History_Age      50000 non-null float64
17 Payment_of_Min_Amount   50000 non-null object
18 Total_EMI_per_month     50000 non-null float64
19 Amount_invested_monthly 47729 non-null object
20 Monthly_Balance         49438 non-null object
dtypes: float64(5), int64(4), object(12)
memory usage: 8.0+ MB

```

Those columns should be numerical values but it has as string now I'm converting them

```

# Assuming our DataFrame is named 'df'
columns_to_convert = ['Age', 'Annual_Income', 'Num_of_Loan', 'Delay_from_due_date', 'Num_of_Delayed_Payment',
                      'Changed_Credit_Limit', 'Outstanding_Debt', 'Amount_invested_monthly', 'Monthly_Balance']

# Convert selected columns to numeric types
df[columns_to_convert] = df[columns_to_convert].apply(pd.to_numeric, errors='coerce')

# Check the data types after conversion
print(df.dtypes)

```

```

Month                    object
Age                     float64
Occupation              object
Annual_Income           float64
Monthly_Inhand_Salary   float64
Num_Bank_Accounts       int64
Num_Credit_Card         int64
Interest_Rate           int64
Num_of_Loan             float64
Delay_from_due_date     int64
Num_of_Delayed_Payment  float64
Changed_Credit_Limit    float64
Num_Credit_Inquiries    float64
Credit_Mix              object
Outstanding_Debt        float64
Credit_Utilization_Ratio float64
Credit_History_Age      float64
Payment_of_Min_Amount   object
Total_EMI_per_month     float64
Amount_invested_monthly float64
Monthly_Balance         float64
dtype: object

```

```
df.isnull().sum()
```

```

Month                    0
Age                     2477
Occupation              0
Annual_Income           3520
Monthly_Inhand_Salary   7498
Num_Bank_Accounts       0
Num_Credit_Card         0
Interest_Rate           0
Num_of_Loan             2436
Delay_from_due_date     0
Num_of_Delayed_Payment  4925
Changed_Credit_Limit    1059
Num_Credit_Inquiries    1035
Credit_Mix              0
Outstanding_Debt        491
Credit_Utilization_Ratio 0
Credit_History_Age      0
Payment_of_Min_Amount   0
Total_EMI_per_month     0
Amount_invested_monthly 4446
Monthly_Balance         568
dtype: int64

```

Work with Monthly_Inhand_Salary columns and fill all missing rows by its mean

```
df['Monthly_Inhand_Salary'].value_counts().tail(45)
```

```
1792.315000    1
1916.121667    1
1966.027083    1
6617.557500    1
3158.531385    1
2714.053557    1
9477.328612    1
7405.809676    1
3303.945000    1
7005.635000    1
873.865417     1
8343.470000    1
6765.796667    1
6204.196083    1
1158.442083    1
8188.211724    1
3750.272500    1
2836.380833    1
8717.732500    1
7823.998705    1
12318.281290   1
2438.323333    1
1801.474167    1
1497.882500    1
926.473333     1
3312.502174    1
14241.603333    1
8437.365531    1
887.077917     1
1488.550000    1
5708.490000    1
4473.288333    1
3792.167500    1
1183.864167    1
11643.130049   1
6107.183333    1
683.253333     1
1039.349084    1
689.187083     1
1652.952966    1
12386.966240   1
5993.870000    1
6763.330000    1
7729.695181    1
2312.785000    1
Name: Monthly_Inhand_Salary, dtype: int64
```

```
average_monthly_salary = df['Monthly_Inhand_Salary'].mean()
df['Monthly_Inhand_Salary'].fillna(value=average_monthly_salary, inplace=True)

df.isnull().sum()
```

```
Month          0
Age           2477
Occupation     0
Annual_Income  3520
Monthly_Inhand_Salary  0
Num_Bank_Accounts  0
Num_Credit_Card  0
Interest_Rate  0
Num_of_Loan    2436
Delay_from_due_date  0
Num_of_Delayed_Payment  4925
Changed_Credit_Limit  1059
Num_Credit_Inquiries  1035
Credit_Mix     0
Outstanding_Debt  491
Credit_Utilization_Ratio  0
Credit_History_Age  0
Payment_of_Min_Amount  0
Total_EMI_per_month  0
Amount_invested_monthly  4446
Monthly_Balance  568
dtype: int64
```

Age never could be negative so that I remove all negative value and fill its by its mean

```
df['Age'].value_counts().head(45)
```

```

39.0      1493
32.0      1440
44.0      1428
22.0      1422
35.0      1414
37.0      1397
27.0      1382
20.0      1374
29.0      1368
26.0      1348
28.0      1344
30.0      1341
41.0      1338
25.0      1325
36.0      1318
24.0      1318
33.0      1280
42.0      1277
19.0      1277
38.0      1266
31.0      1265
21.0      1260
34.0      1236
23.0      1213
45.0      1208
40.0      1196
43.0      1193
46.0      1127
18.0       796
15.0       744
16.0       731
17.0       720
53.0       657
50.0       653
55.0       647
49.0       644
47.0       624
54.0       623
51.0       610
52.0       593
48.0       593
56.0       498
-500.0     464
14.0       181
3287.0      2
Name: Age, dtype: int64

```

```

df['Age'] = df['Age'].apply(lambda x: x if (x >= 0 and x <= 100) else np.nan)
mean_age = df['Age'].mean()
df['Age'].fillna(mean_age, inplace=True)

df.isnull().sum()

```

```

Month      0
Age        0
Occupation 0
Annual_Income      3520
Monthly_Inhand_Salary      0
Num_Bank_Accounts      0
Num_Credit_Card      0
Interest_Rate      0
Num_of_Loan      2436
Delay_from_due_date      0
Num_of_Delayed_Payment      4925
Changed_Credit_Limit      1059
Num_Credit_Inquiries      1035
Credit_Mix      0
Outstanding_Debt      491
Credit_Utilization_Ratio      0
Credit_History_Age      0
Payment_of_Min_Amount      0
Total_EMI_per_month      0
Amount_invested_monthly      4446
Monthly_Balance      568
dtype: int64

```

Working wiht Annual_Income Column

```

df['Annual_Income'].value_counts().head(45)

109945.320      8
17816.750       8
9141.630        8
36585.120       8
72524.200       8

```

```

95596.350      8
22434.160      8
33029.660      7
20867.670      7
40341.160      7
17273.830      6
32543.380      5
43790.400      4
28431.060      4
43268.790      4
108638.760     4
32198.230      4
138920.840     4
56784.540      4
7577.175       4
50807.440      4
13864.835      4
145932.040     4
19183.530      4
68948.320      4
24778.800      4
66189.240      4
20560.130      4
14226.810      4
83552.120      4
37353.580      4
151437.080     4
13000.735      4
14888.915      4
100465.140     4
16697.830      4
20090.020      4
94256.480      4
7295.715       4
18500.540      4
35317.810      4
25703.340      4
70973.320      4
121233.510     4
71518.920      4
Name: Annual_Income, dtype: int64

```

```

# Calculate the mean of the 'Annual_Income' column
mean_annual_income = df['Annual_Income'].mean()

```

```

# Fill missing values with the mean
df['Annual_Income'].fillna(mean_annual_income, inplace=True)
df.isnull().sum()

```

```

Month          0
Age            0
Occupation     0
Annual_Income  0
Monthly_Inhand_Salary  0
Num_Bank_Accounts  0
Num_Credit_Card  0
Interest_Rate  0
Num_of_Loan    2436
Delay_from_due_date  0
Num_of_Delayed_Payment  4925
Changed_Credit_Limit  1059
Num_Credit_Inquiries  1035
Credit_Mix     0
Outstanding_Debt  491
Credit_Utilization_Ratio  0
Credit_History_Age  0
Payment_of_Min_Amount  0
Total_EMI_per_month  0
Amount_invested_monthly  4446
Monthly_Balance  568
dtype: int64

```

```

average_monthly_salary = df['Monthly_Inhand_Salary'].mean()
df['Monthly_Inhand_Salary'].fillna(value=average_monthly_salary, inplace=True)

```

```
df.isnull().sum()
```

```

Month          0
Age            0
Occupation     0
Annual_Income  0
Monthly_Inhand_Salary  0
Num_Bank_Accounts  0
Num_Credit_Card  0
Interest_Rate  0

```



```

Num_of_Loan                2436
Delay_from_due_date        0
Num_of_Delayed_Payment     4925
Changed_Credit_Limit       1059
Num_Credit_Inquiries        1035
Credit_Mix                 0
Outstanding_Debt           491
Credit_Utilization_Ratio   0
Credit_History_Age         0
Payment_of_Min_Amount      0
Total_EMI_per_month        0
Amount_invested_monthly    4446
Monthly_Balance            568
dtype: int64

```

Working wiht Number_of_Delayed_payment columns and fill it's missing rows with it's random vlaues those values are has most frequently in this columns

```
df['Num_of_Delayed_Payment'].value_counts().head(50)
```

```

19.0    2622
15.0    2594
18.0    2570
16.0    2548
17.0    2545
10.0    2517
12.0    2483
11.0    2440
20.0    2422
9.0     2365
8.0     2352
14.0    2007
13.0    2000
21.0    1315
7.0     1140
22.0    1116
6.0     1076
5.0     1036
23.0    1011
3.0      939
4.0      887
2.0      872
24.0     836
1.0      814
25.0     813
0.0      784
26.0     147
-1.0     123
27.0     104
-2.0     103
28.0      64
-3.0      49
2606.0     3
538.0      2
265.0      2
549.0      2
688.0      2
2583.0      2
3064.0      2
1377.0      2
2649.0      2
762.0      2
2608.0      2
861.0      2
3425.0      2
2595.0      2
3666.0      1
1718.0      1
1344.0      1
4206.0      1
Name: Num_of_Delayed_Payment, dtype: int64

```

```

# Create a list of the top most common values
top_values = [19, 15, 18, 16, 17, 10, 11, 12, 20, 9, 8, 14, 13, 21, 7, 22, 6, 5, 23, 3, 4, 2, 24, 1, 25, 0]

# Replace negative numbers with NaN
df['Num_of_Delayed_Payment'] = df['Num_of_Delayed_Payment'].apply(lambda x: x if x >= 0 else np.nan)

# Replace missing values with the top values
missing_indices = df['Num_of_Delayed_Payment'].isnull()
df.loc[missing_indices, 'Num_of_Delayed_Payment'] = np.random.choice(top_values, size=missing_indices.sum())

```

Working with Num_of_Loan columns and fill its missing rows with its random values those values are has most frequently in this columns

```
df.isnull().sum()
```

```

Month                0
Age                  0
Occupation            0
Annual_Income         0
Monthly_Inhand_Salary 0
Num_Bank_Accounts     0
Num_Credit_Card       0
Interest_Rate         0
Num_of_Loan          2436
Delay_from_due_date   0
Num_of_Delayed_Payment 0
Changed_Credit_Limit  1059
Num_Credit_Inquiries  1035
Credit_Mix           0
Outstanding_Debt      491
Credit_Utilization_Ratio 0
Credit_History_Age    0
Payment_of_Min_Amount 0
Total_EMI_per_month   0
Amount_invested_monthly 4446
Monthly_Balance       568
dtype: int64

```

```
df.Num_of_Loan.value_counts().head(25)
```

```

2.0      7173
3.0      7114
4.0      6982
0.0      5163
1.0      5029
6.0      3707
7.0      3483
5.0      3437
-100.0    1974
9.0      1746
8.0      1506
1237.0     2
463.0     2
452.0     2
505.0     2
286.0     2
1475.0     2
1024.0     2
1445.0     2
106.0     2
198.0     2
385.0     2
602.0     2
140.0     2
263.0     2
Name: Num_of_Loan, dtype: int64

```

```
# Replace negative numbers with NaN
```

```
df['Num_of_Loan'] = df['Num_of_Loan'].apply(lambda x: x if x >= 0 else np.nan)
```

```
# Create a list of the top most common values
```

```
top_values = [2.0, 3.0, 4.0, 0.0, 1.0, 6.0, 7.0, 5.0, 9.0, 8.0]
```

```
# Replace missing values with the top values
```

```
missing_indices = df['Num_of_Loan'].isnull()
```

```
df.loc[missing_indices, 'Num_of_Loan'] = np.random.choice(top_values, size=missing_indices.sum())
```

```
df.isnull().sum()
```

```

Month                0
Age                  0
Occupation            0
Annual_Income         0
Monthly_Inhand_Salary 0
Num_Bank_Accounts     0
Num_Credit_Card       0
Interest_Rate         0
Num_of_Loan           0
Delay_from_due_date   0
Num_of_Delayed_Payment 0
Changed_Credit_Limit  1059
Num_Credit_Inquiries  1035
Credit_Mix           0
Outstanding_Debt      491

```

```

Credit_Utilization_Ratio    0
Credit_History_Age          0
Payment_of_Min_Amount       0
Total_EMI_per_month         0
Amount_invested_monthly     4446
Monthly_Balance              568
dtype: int64

```

working with Changed_Credite_Limint column

```

negative_changed_limit = df[df['Changed_Credit_Limit'] < 0]
print(negative_changed_limit['Changed_Credit_Limit'].value_counts().head(10))

```

```

-4.08    5
-0.79    5
-2.50    5
-2.03    5
-1.41    5
-0.60    4
-1.46    4
-3.22    4
-0.67    4
-0.99    4
Name: Changed_Credit_Limit, dtype: int64

```

In a financial context, a negative value for "Changed_Credit_Limit" might represent a reduction or decrease in the credit limit. In some cases, financial institutions or credit card companies may decrease a person's credit limit based on various factors such as changes in creditworthiness, missed payments, or other risk-related considerations.

```

mean_credit_limit = df['Changed_Credit_Limit'].mean()

# Fill missing values with the mean
df['Changed_Credit_Limit'].fillna(mean_credit_limit, inplace=True)
df.isnull().sum()

```

```

Month                0
Age                  0
Occupation           0
Annual_Income        0
Monthly_Inhand_Salary 0
Num_Bank_Accounts    0
Num_Credit_Card      0
Interest_Rate        0
Num_of_Loan          0
Delay_from_due_date  0
Num_of_Delayed_Payment 0
Changed_Credit_Limit 0
Num_Credit_Inquiries 1035
Credit_Mix           0
Outstanding_Debt     491
Credit_Utilization_Ratio 0
Credit_History_Age   0
Payment_of_Min_Amount 0
Total_EMI_per_month   0
Amount_invested_monthly 4446
Monthly_Balance       568
dtype: int64

```

Working with Num_Credit_Inquiries column

```
df['Num_Credit_Inquiries'].value_counts().head(50)
```

```

5.0    4709
4.0    4402
6.0    4375
7.0    4295
8.0    3922
9.0    3523
3.0    3466
11.0   2996
10.0   2982
12.0   2585
2.0    2454
13.0   2065
1.0    1747
14.0   1394
0.0    1102
15.0   1063
16.0    651
17.0    388
2326.0    3

```

```

2338.0      3
1431.0      3
1823.0      3
2019.0      3
1856.0      3
1785.0      3
2179.0      2
1984.0      2
1902.0      2
824.0       2
1990.0      2
2013.0      2
1808.0      2
951.0       2
593.0       2
881.0       2
396.0       2
1879.0      2
695.0       2
1551.0      2
323.0       2
151.0       2
2034.0      2
140.0       2
1138.0      2
2328.0      2
801.0       2
1228.0      2
1694.0      2
2292.0      2
1416.0      2
Name: Num_Credit_Inquiries, dtype: int64

```

```

# Set the number of rows to fill
num_rows_to_fill = 1035

# Generate random values between 1 and 20
random_values = np.random.randint(1, 21, size=num_rows_to_fill)

# Find the indices of missing values in 'Num_Credit_Inquiries'
missing_indices = df['Num_Credit_Inquiries'].isnull()

# Replace missing values with random values
df.loc[missing_indices, 'Num_Credit_Inquiries'] = random_values

df.isnull().sum()

```

```

Month                0
Age                  0
Occupation           0
Annual_Income        0
Monthly_Inhand_Salary 0
Num_Bank_Accounts    0
Num_Credit_Card      0
Interest_Rate        0
Num_of_Loan          0
Delay_from_due_date  0
Num_of_Delayed_Payment 0
Changed_Credit_Limit 0
Num_Credit_Inquiries 0
Credit_Mix           0
Outstanding_Debt     491
Credit_Utilization_Ratio 0
Credit_History_Age   0
Payment_of_Min_Amount 0
Total_EMI_per_month  0
Amount_invested_monthly 4446
Monthly_Balance      568
dtype: int64

```

Working wiht Amount_invested_monthly column

```

df['Amount_invested_monthly'].value_counts().tail(50)

793.597186      1
109.524179      1
82.571977       1
102.658023      1
79.782635       1
285.343947      1
57.534533       1
748.625498      1
49.893509       1
63.272686       1
221.333884      1

```

```

36.952274    1
18.184504    1
13.825182    1
17.913243    1
42.322310    1
158.080629   1
11.908091    1
188.822031   1
445.518708   1
257.173563   1
352.895283   1
188.628242   1
262.440115   1
168.210066   1
77.923151    1
92.694694    1
49.231406    1
40.952555    1
74.746469    1
735.175586   1
130.785144   1
390.765481   1
192.343397   1
409.051633   1
610.385627   1
687.238659   1
793.371509   1
189.034767   1
54.865397    1
41.622649    1
247.478334   1
183.951614   1
147.419348   1
592.143969   1
197.217131   1
366.231484   1
34.899406    1
256.908305   1
220.457878   1
Name: Amount_invested_monthly, dtype: int64

```

```

# Convert 'Amount_invested_monthly' to numeric, handling errors with 'coerce'
#df['Amount_invested_monthly'] = pd.to_numeric(df['Amount_invested_monthly'], errors='coerce')

```

```

# Calculate the mean after converting to numeric
mean_amount_invested = df['Amount_invested_monthly'].mean()

```

```

# Fill missing values with the mean
df['Amount_invested_monthly'].fillna(value=mean_amount_invested, inplace=True)

```

```
df.isnull().sum()
```

```

Month                0
Age                  0
Occupation           0
Annual_Income        0
Monthly_Inhand_Salary 0
Num_Bank_Accounts    0
Num_Credit_Card      0
Interest_Rate        0
Num_of_Loan          0
Delay_from_due_date  0
Num_of_Delayed_Payment 0
Changed_Credit_Limit 0
Num_Credit_Inquiries 0
Credit_Mix           0
Outstanding_Debt     491
Credit_Utilization_Ratio 0
Credit_History_Age  0
Payment_of_Min_Amount 0
Total_EMI_per_month  0
Amount_invested_monthly 0
Monthly_Balance      568
dtype: int64

```

Working wiht Outstanding_Debt column

```
df['Outstanding_Debt'].value_counts().head(23)
```

```

1360.45    12
460.46     12
1109.03    12
1151.70    12
1170.58     8

```

```

2196.59      8
1466.97      8
1286.07      8
2538.06      8
255.82       8
2536.84      8
852.74       8
380.09       8
557.78       8
630.24       8
434.36       8
1334.09      8
1004.26      8
146.68       8
248.84       8
157.62       8
795.69       8
713.33       8
Name: Outstanding_Debt, dtype: int64

```

```

# Calculate the mean of 'Outstanding_Debt' excluding negative values
mean_outstanding_debt = df[df['Outstanding_Debt'] >= 0]['Outstanding_Debt'].mean()

```

```

# Replace negative values with NaN
#df.loc[df['Outstanding_Debt'] < 0, 'Outstanding_Debt'] = np.nan

```

```

# Fill missing values with the mean
df['Outstanding_Debt'].fillna(mean_outstanding_debt, inplace=True)

```

```
df.isnull().sum()
```

```

Month                0
Age                  0
Occupation           0
Annual_Income        0
Monthly_Inhand_Salary 0
Num_Bank_Accounts    0
Num_Credit_Card      0
Interest_Rate        0
Num_of_Loan          0
Delay_from_due_date  0
Num_of_Delayed_Payment 0
Changed_Credit_Limit 0
Num_Credit_Inquiries 0
Credit_Mix          0
Outstanding_Debt     0
Credit_Utilization_Ratio 0
Credit_History_Age  0
Payment_of_Min_Amount 0
Total_EMI_per_month  0
Amount_invested_monthly 0
Monthly_Balance      568
dtype: int64

```

Working with Monthly_Balance column

```
df['Monthly_Balance'].value_counts().sample(50)
```

```

624.496773      1
310.520930      1
308.959343      1
358.643501      1
315.184271      1
252.178460      1
246.547030      1
420.731237      1
41.941683       1
492.359361      1
401.527779      1
241.932909      1
329.325305      1
268.697946      1
948.781690      1
419.795559      1
246.912272      1
245.064532      1
325.761247      1
344.572743      1
425.782800      1
253.403982      1
318.564608      1
515.434613      1
213.966040      1
537.237796      1
365.753263      1
327.169933      1

```

```

747.322558    1
698.819504    1
173.248443    1
351.310832    1
343.605808    1
242.449856    1
123.516632    1
316.542443    1
729.085081    1
456.344600    1
246.815626    1
327.017560    1
336.823280    1
923.144790    1
384.537916    1
344.424166    1
273.264816    1
58.720646     1
211.828301    1
296.721660    1
300.505466    1
616.609056    1
Name: Monthly_Balance, dtype: int64

```

```

mean_monthly_balance = df['Monthly_Balance'].mean()
mean_monthly_balance

```

```
403.0662568630411
```

```

df['Monthly_Balance'].fillna(value=mean_monthly_balance, inplace=True)
df.isnull().sum()

```

```

Month          0
Age            0
Occupation     0
Annual_Income  0
Monthly_Inhand_Salary  0
Num_Bank_Accounts  0
Num_Credit_Card  0
Interest_Rate  0
Num_of_Loan    0
Delay_from_due_date  0
Num_of_Delayed_Payment  0
Changed_Credit_Limit  0
Num_Credit_Inquiries  0
Credit_Mix     0
Outstanding_Debt  0
Credit_Utilization_Ratio  0
Credit_History_Age  0
Payment_of_Min_Amount  0
Total_EMI_per_month  0
Amount_invested_monthly  0
Monthly_Balance  0
dtype: int64

```

```
df.shape
```

```
(50000, 21)
```

I have taken a lot of random values, so that there is high possibility it'll be affected on my result. Cause from every run I would get different results. So that I create a new dataset where has no missing value. And That would be stable and not updated for each run time, It's also help me to get stable result, That I desire most.

```
df.to_csv('new_data.csv', index=False)
```

```

df = pd.read_csv('new_data.csv')
df.sample(20)

```

	Month	Age	Occupation	Annual_Income	Monthly_Inhand_Salary	Num_Bank
33226	November	50.0	Architect	15217.59	999.132500	
20593	October	38.0	Developer	17625.61	1406.800833	
20850	November	52.0	Manager	37367.34	2958.782815	
19638	November	37.0	Teacher	107595.68	8704.306667	
44703	December	20.0	Manager	116897.68	9893.473333	
36028	September	34.0	Musician	8906.93	780.244167	
29539	December	29.0	Teacher	122028.48	9899.040000	
41488	September	27.0	Musician	18525.13	1761.760833	
42655	December	31.0	Musician	62028.36	5048.030000	
10355	December	36.0	Doctor	38186.84	4182.004291	
44508	September	49.0	Lawyer	20770.60	1784.883333	
40385	October	28.0	Writer	20975.39	1572.949167	
9969	October	23.0	Media Manaaer	69865.22	5775.101667	

40007	December	30.0	Engineer	40445.04	4400.004000	
-------	----------	------	----------	----------	-------------	--

EDA & Preprocessing

```
# import ydata_profiling as ydp

# eda_analysis = ydp.ProfileReport(df, title="EDA Analysis of Bank Dataset")

# eda_analysis

# Display data types of each column
data_types = df.dtypes
print(data_types)

# Identify categorical columns
categorical_columns = df.select_dtypes(include=['object']).columns
print("Categorical Columns:", categorical_columns)

Month          object
Age            float64
Occupation     object
Annual_Income  float64
Monthly_Inhand_Salary  float64
Num_Bank_Accounts    int64
Num_Credit_Card      int64
Interest_Rate       int64
Num_of_Loan         float64
Delay_from_due_date  int64
Num_of_Delayed_Payment  float64
Changed_Credit_Limit  float64
Num_Credit_Inquiries  float64
Credit_Mix        object
Outstanding_Debt    float64
Credit_Utilization_Ratio  float64
Credit_History_Age  float64
Payment_of_Min_Amount    object
Total_EMI_per_month    float64
Amount_invested_monthly  float64
Monthly_Balance       float64
dtype: object
Categorical Columns: Index(['Month', 'Occupation', 'Credit_Mix', 'Payment_of_Min_Amount'], dtype='object')

Target Variable is 'Credit_Mix'

Ordinal Variables: 'Payment_of_Min_Amount' is ordinal variable. Since It has a meaningful order, we can use ordinal encoding.

Nominal Variables: 'Month', 'Occupation', 'Payment_Behaviour' are nominal variables. For these, we can use one-hot encoding.
```



```
from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder
```

```
# Define columns for ordinal encoding
ordinal_columns = ['Payment_of_Min_Amount']
ordinal_encoder = OrdinalEncoder()
df[ordinal_columns] = ordinal_encoder.fit_transform(df[ordinal_columns])
df.sample(10)
```


	Month	Age	Occupation	Annual_Income	Monthly_Inhand_Salary	Num_Bank
25984	September	36.0	Accountant	7309.155000	567.096250	
40961	October	23.0	Architect	35970.540000	2728.545000	
18619	December	33.0	Architect	104506.980000	8986.915000	
2336	September	20.0	Teacher	29353.040000	2282.086667	
7914	November	42.0	Media_Manager	69895.760000	5954.646667	
46476	September	27.0	Mechanic	165116.921762	2102.120833	
39932	September	16.0	Doctor	62203.080000	5375.590000	
42994	November	37.0	Accountant	39522.280000	3469.523333	
1083	December	27.0	Developer	34841.870000	4182.004291	
48210	November	21.0	Teacher	34599.800000	4182.004291	

```
df['Month'].value_counts()
```

```
September    12500
October       12500
November      12500
December      12500
Name: Month, dtype: int64
```

```
df['Month'].value_counts().plot(kind='bar')
```

<Axes: >



```
df['Occupation'].value_counts()
```


```
Lawyer        3556
Architect     3436
Mechanic      3423
Engineer      3411
Accountant    3371
Developer     3365
Media_Manager 3356
Scientist     3352
Entrepreneur  3334
Teacher       3322
Journalist    3250
Doctor        3244
Manager       3234
Musician      3188
```

```

Writer          3158
Name: Occupation, dtype: int64

```

```
df['Occupation'].value_counts().plot(kind='bar')
```



```
#df['Payment_Behaviour'].value_counts()
```

✓ Check Outliers

apply Z-score for: Normal Distribution

apply Interquartile-rules for: Skew Distribution

apply Percentile-base(Winsorizing) approche for: Others Distribution

using domain knowledge apply: Box-Cox Transformation and Robust Scaling and log-transformations

```

# Visualize 'Age'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Age', fontsize=16)
sns.histplot(df['Age'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Age')
axes[0].set_xlabel("")
sns.boxplot(x=df['Age'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Age')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()

```

Distribution and Boxplot of Age

Distribution of Age

Boxplot of Age

Removing outliers from Age column

```

upper_limit = df['Age'].mean() + 3 * df['Age'].std()
lower_limit = df['Age'].mean() - 3 * df['Age'].std()


#apply z-score

df['Age'] = np.where(
    df.Age > upper_limit,
    upper_limit,
    np.where(
        df.Age < lower_limit,
        lower_limit,
        df.Age
    )
)

# Visualize 'Age'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Ag after removing outliers', fontsize=16)
sns.histplot(df['Age'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Age after removing outliers')
axes[0].set_xlabel("")
sns.boxplot(x=df['Age'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Age')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()

```

Distribution and Boxplot of Ag after removing outliers




```

# Visualize 'Annual_Income'
pd.set_option('display.float_format', lambda x: '%.3f' % x)
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Annual_Income', fontsize=16)
sns.histplot(df['Annual_Income'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Annual_Income')
axes[0].set_xlabel("")
sns.boxplot(x=df['Annual_Income'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Annual_Income')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()

```

Distribution and Boxplot of Annual_Income



Removing Outliers from Annual Income column

```
upper_limit = df.Annual_Income.quantile(0.99)
lower_limit = df.Annual_Income.quantile(0.01)
```

```
#apply percentile
```

```
df['Annual_Income'] = np.where(
    df.Annual_Income >= upper_limit,
    upper_limit,
    np.where(df.Annual_Income <= lower_limit,
    lower_limit,
    df.Annual_Income)
)
```


```
df['Annual_Income'] = np.log1p(df['Annual_Income']) #apply log-transformations
```

```
df.Annual_Income.describe()
```


```
count    50000.000
mean      10.655
std       0.858
min       8.933
25%       9.907
50%      10.622
75%      11.333
max      12.096
Name: Annual_Income, dtype: float64
```

```
# Visualize 'Annual_Income'
pd.set_option('display.float_format', lambda x: '%.3f' % x)
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Annual_Income after removing outliers', fontsize=16)
sns.histplot(df['Annual_Income'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Annual_Income')
axes[0].set_xlabel("")
sns.boxplot(x=df['Annual_Income'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Annual_Income')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```

Distribution and Boxplot of Annual_Income after removing outliers



```
# Visualize 'Monthly_Inhand_Salary'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Monthly_Inhand_Salary', fontsize=16)
sns.histplot(df['Monthly_Inhand_Salary'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Monthly_Inhand_Salary')
axes[0].set_xlabel("")
sns.boxplot(x=df['Monthly_Inhand_Salary'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Monthly_Inhand_Salary')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```



Removing outliers from Monthly_Inhand_Salary using Interquartile

```
percentile25 = df.Monthly_Inhand_Salary.quantile(0.25)
percentile75 = df.Monthly_Inhand_Salary.quantile(0.75)
```

```
#apply Interquartile
```

```
iqr = percentile75 - percentile25
```

```
upper_limit = percentile75 + 1.5 * iqr
lower_limit = percentile25 - 1.5 * iqr
```


```
df['Monthly_Inhand_Salary'] = np.where(
    df.Monthly_Inhand_Salary > upper_limit,
    upper_limit,
    np.where(
        df.Monthly_Inhand_Salary < lower_limit,
        lower_limit,
        df.Monthly_Inhand_Salary
    )
)
```

```
df.Monthly_Inhand_Salary.describe()
```

```
count    50000.000
mean      4107.302
std       2718.816
min        303.645
25%       1794.304
50%       3848.682
75%       5338.968
max      10655.963
Name: Monthly_Inhand_Salary, dtype: float64
```

```
# Visualize 'Monthly_Inhand_Salary'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Monthly_Inhand_Salary after removing', fontsize=16)
sns.histplot(df['Monthly_Inhand_Salary'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Monthly_Inhand_Salary')
axes[0].set_xlabel("")
sns.boxplot(x=df['Monthly_Inhand_Salary'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Monthly_Inhand_Salary')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```


Distribution and Boxplot of Monthly_Inhand_Salary after removing



Double-click (or enter) to edit

```
# Visualize 'Num_Bank_Accounts'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Num_Bank_Accounts', fontsize=16)
sns.histplot(df['Num_Bank_Accounts'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Num_Bank_Accounts')
axes[0].set_xlabel("")
sns.boxplot(x=df['Num_Bank_Accounts'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Num_Bank_Accounts')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```

Distribution and Boxplot of Num_Bank_Accounts



Apply log Transformer

```
#removing negative values
df['Num_Bank_Accounts'] = df['Num_Bank_Accounts'] + abs(df['Num_Bank_Accounts'].min()) + 1

df.Num_Bank_Accounts.describe()

count    50000.000
mean       18.838
std       116.397
min         1.000
25%         5.000
50%         8.000
75%         9.000
max       1800.000
Name: Num_Bank_Accounts, dtype: float64
```


```
# Apply log transformation
df['Num_Bank_Accounts'] = np.log1p(df['Num_Bank_Accounts'])
```

```
# Display summary statistics
print(df['Num_Bank_Accounts'].describe())
```

```
count    50000.000
mean       2.123
std        0.624
min        0.693
25%        1.792
50%        2.197
75%        2.303
max        7.496
Name: Num_Bank_Accounts, dtype: float64
```

```
# Visualize 'Num_Bank_Accounts'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Num_Bank_Accounts After Removing outliers', fontsize=16)
sns.histplot(df['Num_Bank_Accounts'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Num_Bank_Accounts')
axes[0].set_xlabel("")
sns.boxplot(x=df['Num_Bank_Accounts'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Num_Bank_Accounts')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```

Distribution and Boxplot of Num_Bank_Accounts After Removing outliers



```
# Visualize 'Num_Credit_Card'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Num_Credit_Card', fontsize=16)
sns.histplot(df['Num_Credit_Card'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Num_Credit_Card')
axes[0].set_xlabel("")
sns.boxplot(x=df['Num_Credit_Card'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Num_Credit_Card')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```

Distribution and Boxplot of Num_Credit_Card



Num_Credit_Card for this column I don't understand it's distribution that's why I'm using Box-Cox transformation


```
#removing negative values
df['Num_Credit_Card'] = df['Num_Credit_Card'] + abs(df['Num_Credit_Card'].min()) + 1
df['Num_Bank_Accounts'] = np.log1p(df['Num_Bank_Accounts'])
```

```
df.Num_Credit_Card.describe()
```

```
count    50000.000
mean       23.921
std       129.315
min         1.000
25%         5.000
50%         6.000
75%         8.000
max       1500.000
Name: Num_Credit_Card, dtype: float64
```

```
# Visualize 'Interest_Rate'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Interest_Rate', fontsize=16)
sns.histplot(df['Interest_Rate'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Interest_Rate')
axes[0].set_xlabel("")
sns.boxplot(x=df['Interest_Rate'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Interest_Rate')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```

Distribution and Boxplot of Interest_Rate



```
from scipy.stats import boxcox
```

```
transformed_data, lambda_value = boxcox(df['Interest_Rate'], lmbda=None)
```

```
df['Interest_Rate'] = np.power((transformed_data * lambda_value) + 1, 1 / lambda_value)
```


```
df.Interest_Rate.describe()
```

```
count    50000.000
mean       68.773
std       451.602
min         1.000
25%         8.000
50%        13.000
75%        20.000
max       5799.000
Name: Interest_Rate, dtype: float64
```




```
# Visualize 'Interest_Rate'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Interest_Rate after removing outliers', fontsize=16)
sns.histplot(df['Interest_Rate'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Interest_Rate')
axes[0].set_xlabel("")
sns.boxplot(x=df['Interest_Rate'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Interest_Rate')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```

Distribution and Boxplot of Interest_Rate after removing outliers



```
# Visualize 'Num_of_Loan'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Num_of_Loan', fontsize=16)
sns.histplot(df['Num_of_Loan'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Num_of_Loan')
axes[0].set_xlabel("")
sns.boxplot(x=df['Num_of_Loan'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Num_of_Loan')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```

Distribution and Boxplot of Num_of_Loan



```
#removing negative values
df['Num_of_Loan'] = df['Num_of_Loan'] + abs(df['Num_of_Loan'].min()) + 1

df['Num_of_Loan'] = np.log1p(df['Num_of_Loan'])

df.Num_of_Loan.describe()
```

```
count    50000.000
mean         1.640
std         0.595
min         0.693
```

```

25%      1.386
50%      1.609
75%      1.946
max       7.312
Name: Num_of_Loan, dtype: float64


```

```

# Visualize 'Num_of_Loan'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Num_of_Loan after removing outliers', fontsize=16)
sns.histplot(df['Num_of_Loan'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Num_of_Loan')
axes[0].set_xlabel("")
sns.boxplot(x=df['Num_of_Loan'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Num_of_Loan')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()

```

Distribution and Boxplot of Num_of_Loan after removing outliers




```

# Visualize 'Delay_from_due_date'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Delay_from_due_date', fontsize=16)
sns.histplot(df['Delay_from_due_date'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Delay_from_due_date')
axes[0].set_xlabel("")
sns.boxplot(x=df['Delay_from_due_date'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Delay_from_due_date')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()

```

Distribution and Boxplot of Delay_from_due_date



Removing outliers from Delay_from_due_date column

```

percentile25 = df.Delay_from_due_date.quantile(0.25)
percentile75 = df.Delay_from_due_date.quantile(0.75)


iqr = percentile75 - percentile25

upper_limit = percentile75 + 1.5 * iqr
lower_limit = percentile25 - 1.5 * iqr

df['Delay_from_due_date'] = np.where(
    df.Delay_from_due_date > upper_limit,
    upper_limit,
    np.where(
        df.Delay_from_due_date < lower_limit,
        lower_limit,
        df.Delay_from_due_date
    )
)

# Visualize 'Delay_from_due_date'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Delay_from_due_date after removing outliers', fontsize=16)
sns.histplot(df['Delay_from_due_date'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Delay_from_due_date')
axes[0].set_xlabel("")
sns.boxplot(x=df['Delay_from_due_date'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Delay_from_due_date')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()

```



```

# Visualize 'Num_Credit_Card'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Num_Credit_Card', fontsize=16)
sns.histplot(df['Num_Credit_Card'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Num_Credit_Card')
axes[0].set_xlabel("")
sns.boxplot(x=df['Num_Credit_Card'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Num_Credit_Card')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()

```

Distribution and Boxplot of Num_Credit_Card




```
df['Num_Credit_Card'] = np.log1p(df['Num_Credit_Card'])
```

```
df.Num_Credit_Card.describe()
```

```
count    50000.000
mean       2.082
std        0.735
min        0.693
25%        1.792
50%        1.946
75%        2.197
max        7.314
Name: Num_Credit_Card, dtype: float64
```


```
# Visualize 'Num_Credit_Card'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Num_Credit_Card after removing outliers', fontsize=16)
sns.histplot(df['Num_Credit_Card'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Num_Credit_Card')
axes[0].set_xlabel("")
sns.boxplot(x=df['Num_Credit_Card'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Num_Credit_Card')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```

Distribution and Boxplot of Num_Credit_Card after removing outliers



```
# Visualize 'Num_Credit_Inquiries'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Num_Credit_Inquiries', fontsize=16)
sns.histplot(df['Num_Credit_Inquiries'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Num_Credit_Inquiries')
axes[0].set_xlabel("")
sns.boxplot(x=df['Num_Credit_Inquiries'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Num_Credit_Inquiries')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```

Distribution and Boxplot of Num_Credit_Inquiries



```
from scipy.stats import boxcox
```

```
#removing negative values
```

```
df['Num_Credit_Inquiries'] = df['Num_Credit_Inquiries'] + abs(df['Num_Credit_Inquiries'].min()) + 1
```

```
transformed_data, lambda_value = boxcox(df['Num_Credit_Inquiries'], lmbda=None)
```

```
df['Num_Credit_Inquiries'] = np.power((transformed_data * lambda_value) + 1, 1 / lambda_value)
```

```
df.Num_Credit_Inquiries.describe()
```

```
count    50000.000
mean       30.685
std       194.955
min         1.000
25%         5.000
50%         8.000
75%        12.000
max       2594.000
Name: Num_Credit_Inquiries, dtype: float64
```

```
# Visualize 'Num_Credit_Inquiries'
```

```
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
```

```
fig.suptitle('Distribution and Boxplot of Num_Credit_Inquiries after removing outliers', fontsize=16)
```

```
sns.histplot(df['Num_Credit_Inquiries'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
```

```
axes[0].set_title('Distribution of Num_Credit_Inquiries')
```

```
axes[0].set_xlabel("")
```


```
sns.boxplot(x=df['Num_Credit_Inquiries'], ax=axes[1], color='salmon')
```

```
axes[1].set_title('Boxplot of Num_Credit_Inquiries')
```

```
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
```

```
plt.show()
```

Distribution and Boxplot of Num_Credit_Inquiries after removing outliers



```
# Visualize 'Outstanding_Debt'
```

```
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
```

```
fig.suptitle('Distribution and Boxplot of Outstanding_Debt', fontsize=16)
```

```
sns.histplot(df['Outstanding_Debt'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
```

```
axes[0].set_title('Distribution of Outstanding_Debt')
```

```
axes[0].set_xlabel("")
```


```
sns.boxplot(x=df['Outstanding_Debt'], ax=axes[1], color='salmon')
```

```
axes[1].set_title('Boxplot of Outstanding_Debt')
```

```
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
```

```
plt.show()
```

Distribution and Boxplot of Outstanding_Debt



Removing outliers from Outstanding_Debt column

```
percentile25 = df.Outstanding_Debt.quantile(0.25)
percentile75 = df.Outstanding_Debt.quantile(0.75)


iqr = percentile75 - percentile25

upper_limit = percentile75 + 1.5 * iqr
lower_limit = percentile25 - 1.5 * iqr

df['Delay_from_due_date'] = np.where(
    df.Outstanding_Debt > upper_limit,
    upper_limit,
    np.where(
        df.Outstanding_Debt < lower_limit,
        lower_limit,
        df.Outstanding_Debt
    )
)


# Visualize 'Outstanding_Debt'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Outstanding_Debt', fontsize=16)
sns.histplot(df['Outstanding_Debt'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Outstanding_Debt')
axes[0].set_xlabel("")
sns.boxplot(x=df['Outstanding_Debt'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Outstanding_Debt')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```

Distribution and Boxplot of Outstanding_Debt




```
# Visualize 'Credit_Utilization_Ratio'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Credit_Utilization_Ratio', fontsize=16)
sns.histplot(df['Credit_Utilization_Ratio'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Credit_Utilization_Ratio')
axes[0].set_xlabel("")
sns.boxplot(x=df['Credit_Utilization_Ratio'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Credit_Utilization_Ratio')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```

Distribution and Boxplot of Credit_Utilization_Ratio




```
# Visualize 'Credit_History_Age'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Credit_History_Age', fontsize=16)
sns.histplot(df['Credit_History_Age'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Credit_History_Age')
axes[0].set_xlabel("")
sns.boxplot(x=df['Credit_History_Age'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Credit_History_Age')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```

Distribution and Boxplot of Credit_History_Age




```
# Visualize 'Payment_of_Min_Amount'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Payment_of_Min_Amount', fontsize=16)
sns.histplot(df['Payment_of_Min_Amount'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Payment_of_Min_Amount')
axes[0].set_xlabel("")
sns.boxplot(x=df['Payment_of_Min_Amount'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Payment_of_Min_Amount')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```

Distribution and Boxplot of Payment_of_Min_Amount



```
# Visualize 'Total_EMI_per_month'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Total_EMI_per_month', fontsize=16)
sns.histplot(df['Total_EMI_per_month'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Total_EMI_per_month')
axes[0].set_xlabel("")
sns.boxplot(x=df['Total_EMI_per_month'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Total_EMI_per_month')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```

Distribution and Boxplot of Total_EMI_per_month




```
df['Total_EMI_per_month'] = np.log1p(df['Total_EMI_per_month'])
```

```
df.Total_EMI_per_month.describe()
```

```
count    50000.000
mean       4.207
std        1.980
min         0.000
25%        3.503
50%        4.327
75%        5.177
max       11.319
Name: Total_EMI_per_month, dtype: float64
```


```
# Visualize 'Total_EMI_per_month'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Total_EMI_per_month after removing outliers', fontsize=16)
sns.histplot(df['Total_EMI_per_month'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Total_EMI_per_month')
axes[0].set_xlabel("")
sns.boxplot(x=df['Total_EMI_per_month'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Total_EMI_per_month')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```


Distribution and Boxplot of Total_EMI_per_month after removing outliers



```
# Visualize 'Amount_invested_monthly'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Amount_invested_monthly', fontsize=16)
sns.histplot(df['Amount_invested_monthly'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Amount_invested_monthly')
axes[0].set_xlabel("")
sns.boxplot(x=df['Amount_invested_monthly'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Amount_invested_monthly')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```

Distribution and Boxplot of Amount_invested_monthly



```
percentile25 = df.Amount_invested_monthly.quantile(0.25)
percentile75 = df.Amount_invested_monthly.quantile(0.75)
```


```
iqr = percentile75 - percentile25
```

```
upper_limit = percentile75 + 1.5 * iqr
lower_limit = percentile25 - 1.5 * iqr
```

```
df['Amount_invested_monthly'] = np.where(
    df.Amount_invested_monthly > upper_limit,
    upper_limit,
    np.where(
        df.Amount_invested_monthly < lower_limit,
        lower_limit,
        df.Amount_invested_monthly
    )
)
```


```
# Visualize 'Amount_invested_monthly'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Amount_invested_monthly after removing outliers', fontsize=16)
sns.histplot(df['Amount_invested_monthly'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Amount_invested_monthly')
axes[0].set_xlabel("")
sns.boxplot(x=df['Amount_invested_monthly'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Amount_invested_monthly')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```

Distribution and Boxplot of Amount_invested_monthly after removing outliers



```
# Visualize 'Monthly_Balance'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
fig.suptitle('Distribution and Boxplot of Monthly_Balance', fontsize=16)
sns.histplot(df['Monthly_Balance'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
axes[0].set_title('Distribution of Monthly_Balance')
axes[0].set_xlabel("")
sns.boxplot(x=df['Monthly_Balance'], ax=axes[1], color='salmon')
axes[1].set_title('Boxplot of Monthly_Balance')
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
plt.show()
```

Distribution and Boxplot of Monthly_Balance



```
percentile25 = df.Monthly_Balance.quantile(0.25)
percentile75 = df.Monthly_Balance.quantile(0.75)
```

```
iqr = percentile75 - percentile25
```

```
upper_limit = percentile75 + 1.5 * iqr
lower_limit = percentile25 - 1.5 * iqr
```

```
df['Monthly_Balance'] = np.where(
    df.Monthly_Balance > upper_limit,
    upper_limit,
    np.where(
        df.Monthly_Balance < lower_limit,
        lower_limit,
        df.Monthly_Balance
    )
)
```

```
df.Monthly_Balance.describe()
```

```
count    50000.000
mean      388.694
std       173.506
min        0.103
25%       271.119
50%       338.992
75%       468.571
max       764.750
Name: Monthly_Balance, dtype: float64
```

```
# Visualize 'Monthly_Balance'
```

```
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 5))
```

```
fig.suptitle('Distribution and Boxplot of Monthly_Balance after removing outliers', fontsize=16)
```

```
sns.histplot(df['Monthly_Balance'], ax=axes[0], kde=True, color='skyblue', stat="density", common_norm=False)
```

```
axes[0].set_title('Distribution of Monthly_Balance')
```

```
axes[0].set_xlabel("")
```


```
sns.boxplot(x=df['Monthly_Balance'], ax=axes[1], color='salmon')
```

```
axes[1].set_title('Boxplot of Monthly_Balance')
```

```
plt.tight_layout(rect=[0, 0.03, 1, 0.9])
```

```
plt.show()
```

Distribution and Boxplot of Monthly_Balance after removing outliers



✓ Choosing Target Variable

"Monthly Balance" can have different meanings depending on the context, but here are the two most likely interpretations:

- 1. Financial Account Balance:** This is the most common meaning. It refers to the amount of money remaining in your bank account at the end of each month. This applies to checking accounts, savings accounts, credit card accounts, or any other type of financial account.
- 2. Budget Tracking:** In budgeting, "Monthly Balance" refers to the money left after paying all your bills and expenses for the month. This helps track progress towards financial goals like saving for a house or paying off debt.

```
X = df.drop('Credit_Mix', axis=1)
y = df['Credit_Mix']
```

```
y.value_counts()

Standard    18379
Good       12260
neutral     9805
Bad         9556
Name: Credit_Mix, dtype: int64
```

Imbalanced data Handle

```
standard = (18379/(18379+12260+9805+9556)) *100
good = (12260/(18379+12260+9805+9556)) *100
neutral = (9805/(18379+12260+9805+9556)) *100
bad = (9556/(18379+12260+9805+9556)) *100

print('Standard {} percent of total Credit Mix'.format(standard))
print('Good {} percent of total Credit Mix'.format(good))
print('Neutral {} percent of total Credit Mix'.format(neutral))
print('Bad {} percent of total Credit Mix'.format(bad))
```

```
Standard 36.758 percent of total Credit Mix
Good 24.52 percent of total Credit Mix
Neutral 19.61 percent of total Credit Mix
Bad 19.11200000000002 percent of total Credit Mix
```

```
from imblearn.over_sampling import RandomOverSampler
```

```
over = RandomOverSampler(random_state=100)
```

```
X_new , y_new = over.fit_resample(X,y)
```

```
X_new.shape, y_new.shape
```

```
((73516, 20), (73516,))
```

```
y_new.value_counts()
```

```
Good        18379
neutral     18379
Standard    18379
Bad         18379
Name: Credit_Mix, dtype: int64
```

```
X_new.head()
```

	Month	Age	Occupation	Annual_Income	Monthly_Inhand_Salary	Num_Bank_Accoi
0	September	23.000	Scientist	9.858	1824.843	1
1	October	24.000	Scientist	9.858	1824.843	1
2	November	24.000	Scientist	9.858	1824.843	1
3	December	33.820	Scientist	9.858	4182.004	1
4	September	28.000	Teacher	10.459	3037.987	0

Apply MinMax() for Numerical Columns and oneHOT of categorical columns and then apply PCA

```
columns_to_onehot = ['Month', 'Occupation'] # 'Payment_Behaviour'
```

```
# Apply one-hot encoding to the specified columns
X_new = pd.get_dummies(X_new, columns=columns_to_onehot)
X_new.sample(10)
```

	Age	Annual_Income	Monthly_Inhand_Salary	Num_Bank_Accounts	Num_Credit_Ca
41738	32.000	9.619	971.158	1.162	1.7
12641	20.000	12.014	1770.714	1.271	2.4
6773	39.000	12.014	2745.302	0.741	1.7
47575	40.000	12.014	1616.400	1.162	1.7

Split Dataset

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X_new, y_new, test_size=0.2, random_state=42)

X_train.shape, X_test.shape, y_train.shape, y_test.shape

((58812, 37), (14704, 37), (58812,), (14704,))
```

Apply MinMax

```
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

X_train_scaled_df = pd.DataFrame(X_train_scaled)
X_test_scaled_df = pd.DataFrame(X_test_scaled)

X_train_scaled_df.head()

   0    1    2    3    4    5    6    7    8    9   10   11
0  0.422  0.800  1.003  0.413  0.140 -0.091 -0.103 -0.135 -0.091 -0.371 -0.081 -0.155
1  0.227 -1.051 -0.956 -0.012 -0.416 -0.113  0.888 -1.097 -0.101  0.691 -0.107 -1.051
2  2.172 -1.429  0.044 -0.012 -0.665 -0.131 -0.964 -0.444 -0.111 -1.226 -0.142 -0.443
3  0.033  0.802  1.067  0.216  0.414 -0.093  1.066  2.142 -0.050  2.496 -0.101  2.411
4  1.783 -0.265 -0.370 -0.281  0.140 -0.147 -1.647 -1.268 -0.131 -0.495 -0.142 -1.209

X_train_scaled_df.shape, X_test_scaled_df.shape

((58812, 37), (14704, 37))
```

Apply PCA

```
from sklearn.decomposition import PCA

pca = PCA(n_components=3) # Adjust the number of components as needed

X_train_pca = pca.fit_transform(X_train_scaled_df)
X_test_pca = pca.transform(X_test_scaled_df)

X_train_pca_df = pd.DataFrame(X_train_pca, columns=[f'pca_{i+1}' for i in range(3)])
X_test_pca_df = pd.DataFrame(X_test_pca, columns=[f'pca_{i+1}' for i in range(3)])

X_train_pca_df.head()
```

	pca_1	pca_2	pca_3
0	-0.277	1.132	-0.185

```
X_test_pca_df.head()
```

	pca_1	pca_2	pca_3
0	3.000	-1.550	1.454
1	3.984	-0.554	-1.193
2	-0.363	-1.472	1.315
3	1.730	0.017	-1.219
4	2.322	0.536	1.477

```
X_train_pca_df.shape
```

```
(58812, 3)
```


```
import plotly.express as px
```

```
fig = px.scatter_3d(X_train_pca_df, x='pca_1', y='pca_2', z='pca_3', color=y_train)
```

```
fig.update_layout(scene=dict(
    xaxis_title='PC1',
    yaxis_title='PC2',
    zaxis_title='PC3'),
    title='3D Visualization of Training Data after PCA')
```

```
fig.show()
```

3D Visualization of Training Data after PCA




```
import plotly.express as px
```

```
fig = px.scatter_3d(X_test_pca_df, x='pca_1', y='pca_2', z='pca_3', color=y_test)
```

```
fig.update_layout(
    scene=dict(
        xaxis_title='PC1',
        yaxis_title='PC2',
        zaxis_title='PC3'
    ),
    title='3D Visualization of Test Data after PCA'
)
```

```
fig.show()
```

3D Visualization of Test Data after PCA



3. Model Selection:

Logistic Regression

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix
```

```
# Initialize logistic regression model
logreg = LogisticRegression(random_state=42)
```

```
# Fit the model on the training data
logreg.fit(X_train_pca_df, y_train)
```

```
# Make predictions on the testing data
y_pred = logreg.predict(X_test_pca_df)
```

```
# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
f1 = f1_score(y_test, y_pred, average='weighted')
conf_matrix = confusion_matrix(y_test, y_pred)
```

```
print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1 Score: {f1:.4f}")
```

```
print("Confusion Matrix:")
print(conf_matrix)
```

```
Accuracy: 0.5195
Precision: 0.4812
Recall: 0.5195
F1 Score: 0.4876
Confusion Matrix:
[[3032    0  177  573]
 [  11 2907  496  213]
 [ 582  966 1022 1109]
 [ 967 1299  673  677]]
```

Apply RandomForestClassifier

```

from sklearn.ensemble import RandomForestClassifier

# Initialize Random Forest Classifier
rf_classifier = RandomForestClassifier(random_state=42)

# Fit the model on the training data
rf_classifier.fit(X_train_pca_df, y_train)

# Make predictions on the testing data
y_pred = rf_classifier.predict(X_test_pca_df)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
f1 = f1_score(y_test, y_pred, average='weighted')
conf_matrix = confusion_matrix(y_test, y_pred)

# Print the evaluation metrics
print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1 Score: {f1:.4f}")

# Display the confusion matrix
print("Confusion Matrix:")
print(conf_matrix)

```

```

Accuracy: 0.7121
Precision: 0.7091
Recall: 0.7121
F1 Score: 0.7045
Confusion Matrix:
[[3357   1  318 106]
 [   8 3060  367 192]
 [  511  874 1879 415]
 [  408  528  505 2175]]

```

Support Vector Machine Classifier

```

from sklearn.svm import SVC

# Initialize Support Vector Machine Classifier
svm_classifier = SVC(random_state=42)

# Fit the model on the training data
svm_classifier.fit(X_train_pca_df, y_train)

# Make predictions on the testing data
y_pred_svm = svm_classifier.predict(X_test_pca_df)

# Evaluate the model
accuracy_svm = accuracy_score(y_test, y_pred_svm)
precision_svm = precision_score(y_test, y_pred_svm, average='weighted')
recall_svm = recall_score(y_test, y_pred_svm, average='weighted')
f1_svm = f1_score(y_test, y_pred_svm, average='weighted')
conf_matrix_svm = confusion_matrix(y_test, y_pred_svm)

# Print the evaluation metrics for SVM
print("Support Vector Machine (SVM) Metrics:")
print(f"Accuracy: {accuracy_svm:.4f}")
print(f"Precision: {precision_svm:.4f}")
print(f"Recall: {recall_svm:.4f}")
print(f"F1 Score: {f1_svm:.4f}")

# Display the confusion matrix for SVM
print("Confusion Matrix:")
print(conf_matrix_svm)

```

```

Support Vector Machine (SVM) Metrics:
Accuracy: 0.5592
Precision: 0.4153
Recall: 0.5592
F1 Score: 0.4693
Confusion Matrix:
[[3404   1  377   0]
 [  18 3247  362   0]

```



```
[ 885 1222 1572   0]
[1211 1512   893   0]]
```

Gradient Boosting Classifier (e.g., XGBoost)

Logistic Regression, Random Forest Classifier, and SVM often handle categorical labels directly without the need for encoding.

XGBoost requires label encoding for categorical labels.

```
from sklearn.preprocessing import LabelEncoder
from xgboost import XGBClassifier

# Convert class labels to integers
label_encoder = LabelEncoder()
y_train_encoded = label_encoder.fit_transform(y_train)

# Initialize XGBoost Classifier
xgb_classifier = XGBClassifier(random_state=42)

# Fit the model on the training data
xgb_classifier.fit(X_train_pca_df, y_train_encoded)

# Make predictions on the testing data
y_test_encoded = label_encoder.transform(y_test)
y_pred_xgb = xgb_classifier.predict(X_test_pca_df)

# Evaluate the model
accuracy_xgb = accuracy_score(y_test_encoded, y_pred_xgb)
precision_xgb = precision_score(y_test_encoded, y_pred_xgb, average='weighted')
recall_xgb = recall_score(y_test_encoded, y_pred_xgb, average='weighted')
f1_xgb = f1_score(y_test_encoded, y_pred_xgb, average='weighted')
conf_matrix_xgb = confusion_matrix(y_test_encoded, y_pred_xgb)

# Print the evaluation metrics for XGBoost
print("XGBoost Metrics:")
print(f"Accuracy: {accuracy_xgb:.4f}")
print(f"Precision: {precision_xgb:.4f}")
print(f"Recall: {recall_xgb:.4f}")
print(f"F1 Score: {f1_xgb:.4f}")

# Display the confusion matrix for XGBoost
print("Confusion Matrix:")
print(conf_matrix_xgb)
```

```
XGBoost Metrics:
Accuracy: 0.5700
Precision: 0.5424
Recall: 0.5700
F1 Score: 0.5126
Confusion Matrix:
[[3300    1  409    72]
 [   13 3063   435  116]
 [   763 1045 1709   162]
 [ 1020 1276 1010   310]]
```

I will apply Hyperparameter Tuning on RandomForestClassifier and XGBClassifier Because I got best 'Accuracy' from them.

✓ 5. Hyperparameter Tuning:

Grid Search for Gradient Boosting Classifier (e.g., XGBoost)

```

from sklearn.model_selection import GridSearchCV

# Assuming X_train_pca_df, X_test_pca_df, y_train_encoded, y_test_encoded are defined

# Initialize XGBoost Classifier
xgb_classifier = XGBClassifier(random_state=42)

# Define the parameter grid for Grid Search
param_grid = {
    'n_estimators': [200],
    'max_depth': [7],
    'learning_rate': [0.2],
    'subsample': [0.8],
    'colsample_bytree': [1.0],
    'min_child_weight': [1, 3, 5]
}

# Initialize Grid Search
grid_search = GridSearchCV(xgb_classifier, param_grid, cv=3, scoring='accuracy', n_jobs=-1)

# Fit the Grid Search on the training data
grid_search.fit(X_train_pca_df, y_train_encoded)

# Print the best parameters found by Grid Search
print("Best Parameters:", grid_search.best_params_)

# Make predictions on the testing data using the best model
y_pred_grid_search = grid_search.best_estimator_.predict(X_test_pca_df)

# Evaluate the model after Grid Search
accuracy_grid_search = accuracy_score(y_test_encoded, y_pred_grid_search)
precision_grid_search = precision_score(y_test_encoded, y_pred_grid_search, average='weighted')
recall_grid_search = recall_score(y_test_encoded, y_pred_grid_search, average='weighted')
f1_grid_search = f1_score(y_test_encoded, y_pred_grid_search, average='weighted')
conf_matrix_grid_search = confusion_matrix(y_test_encoded, y_pred_grid_search)

# Print the evaluation metrics after Grid Search
print("Metrics after Grid Search:")
print(f"Accuracy: {accuracy_grid_search:.4f}")
print(f"Precision: {precision_grid_search:.4f}")
print(f"Recall: {recall_grid_search:.4f}")
print(f"F1 Score: {f1_grid_search:.4f}")
print("Confusion Matrix:")
print(conf_matrix_grid_search)

Best Parameters: {'colsample_bytree': 1.0, 'learning_rate': 0.2, 'max_depth': 7, 'min_child_weight': 1, 'n_estimators': 200, 'subsan
Metrics after Grid Search:
Accuracy: 0.5860
Precision: 0.5679
Recall: 0.5860
F1 Score: 0.5535
Confusion Matrix:
[[3191   2  458  131]
 [  14 2942  439  232]
 [ 645  942 1768  324]
 [ 842 1094  965  715]]

```

Random Search for Gradient Boosting Classifier (e.g., XGBoost)

```

from sklearn.model_selection import RandomizedSearchCV

# Assuming X_train_pca_df, X_test_pca_df, y_train_encoded, y_test_encoded are defined

# Initialize XGBoost Classifier
xgb_classifier = XGBClassifier(random_state=42)

# Define the parameter distribution for Random Search
param_dist = {
    'n_estimators': [100],
    'max_depth': [7],
    'learning_rate': [0.2],
    'subsample': [0.8],
    'colsample_bytree': [1.0],
    'min_child_weight': [1],
    'gamma': [0.2],
    'reg_alpha': [0],
    'reg_lambda': [1.0]
}

```

```

# Initialize Random Search
random_search = RandomizedSearchCV(xgb_classifier, param_distributions=param_dist, n_iter=10, cv=3, scoring='accuracy', n_jobs=-1, random

# Fit the Random Search on the training data
random_search.fit(X_train_pca_df, y_train_encoded)

# Print the best parameters found by Random Search
print("Best Parameters:", random_search.best_params_)

# Make predictions on the testing data using the best model
y_pred_random_search = random_search.best_estimator_.predict(X_test_pca_df)

# Evaluate the model after Random Search
accuracy_random_search = accuracy_score(y_test_encoded, y_pred_random_search)
precision_random_search = precision_score(y_test_encoded, y_pred_random_search, average='weighted')
recall_random_search = recall_score(y_test_encoded, y_pred_random_search, average='weighted')
f1_random_search = f1_score(y_test_encoded, y_pred_random_search, average='weighted')
conf_matrix_random_search = confusion_matrix(y_test_encoded, y_pred_random_search)

# Print the evaluation metrics after Random Search
print("Metrics after Random Search:")
print(f"Accuracy: {accuracy_random_search:.4f}")
print(f"Precision: {precision_random_search:.4f}")
print(f"Recall: {recall_random_search:.4f}")
print(f"F1 Score: {f1_random_search:.4f}")
print("Confusion Matrix:")
print(conf_matrix_random_search)

```

```

Best Parameters: {'subsample': 0.8, 'reg_lambda': 1.0, 'reg_alpha': 0, 'n_estimators': 100, 'min_child_weight': 1, 'max_depth': 7,
Metrics after Random Search:
Accuracy: 0.5713
Precision: 0.5502
Recall: 0.5713
F1 Score: 0.5143
Confusion Matrix:
[[3279    0  432   71]
 [  15 3085  430   97]
 [ 770 1044 1713  152]
[1009 1272 1012  323]]

```

Random Forest Classifier: Grid Search:

```

from sklearn.model_selection import GridSearchCV

# Initialize Random Forest Classifier
rf_classifier = RandomForestClassifier(random_state=42)

# Define the parameter grid for Grid Search
param_grid_rf = {
    'n_estimators': [100],
    'max_depth': [None],
    'min_samples_split': [2],
    'min_samples_leaf': [1],
    'bootstrap': [True],
    'max_features': ['auto'],
    'criterion': ['entropy']
}

# Initialize Grid Search
grid_search_rf = GridSearchCV(rf_classifier, param_grid_rf, cv=3, scoring='accuracy', n_jobs=-1)

# Fit the Grid Search on the training data
grid_search_rf.fit(X_train_pca_df, y_train)

# Print the best parameters found by Grid Search
print("Best Parameters (Random Forest):", grid_search_rf.best_params_)

# Make predictions on the testing data using the best model
y_pred_grid_search_rf = grid_search_rf.best_estimator_.predict(X_test_pca_df)

# Evaluate the model after Grid Search
accuracy_grid_search_rf = accuracy_score(y_test, y_pred_grid_search_rf)
precision_grid_search_rf = precision_score(y_test, y_pred_grid_search_rf, average='weighted')
recall_grid_search_rf = recall_score(y_test, y_pred_grid_search_rf, average='weighted')
f1_grid_search_rf = f1_score(y_test, y_pred_grid_search_rf, average='weighted')
conf_matrix_grid_search_rf = confusion_matrix(y_test, y_pred_grid_search_rf)

# Print the evaluation metrics after Grid Search
print("Metrics after Grid Search (Random Forest):")
print(f"Accuracy: {accuracy_grid_search_rf:.4f}")
print(f"Precision: {precision_grid_search_rf:.4f}")
print(f"Recall: {recall_grid_search_rf:.4f}")
print(f"F1 Score: {f1_grid_search_rf:.4f}")
print("Confusion Matrix:")
print(conf_matrix_grid_search_rf)

Best Parameters (Random Forest): {'bootstrap': True, 'criterion': 'entropy', 'max_depth': None, 'max_features': 'auto', 'min_samples':
Metrics after Grid Search (Random Forest):
Accuracy: 0.7145
Precision: 0.7119
Recall: 0.7145
F1 Score: 0.7073
Confusion Matrix:
[[3361    0  323   98]
 [   8 3052  371  196]
 [ 490  871 1910  408]
 [ 416  523  494 2183]]

```

Random Forest Classifier: Random Search:

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix
import numpy as np

# Assuming X_train_pca_df, X_test_pca_df, y_train, y_test are defined

# Initialize Random Forest Classifier
rf_classifier = RandomForestClassifier(random_state=42)

# Define the parameter distribution for Random Search
param_dist_rf = {
    'n_estimators': [100],
    'max_depth': [20],
    'min_samples_split': [5],
    'min_samples_leaf': [4],
    'bootstrap': [True],
    'max_features': ['log2'],
    'criterion': ['entropy'],
    'max_samples': [0.9],
    'min_weight_fraction_leaf': [0.0]
}

# Initialize Random Search
random_search_rf = RandomizedSearchCV(rf_classifier, param_distributions=param_dist_rf, n_iter=10, cv=3, scoring='accuracy', n_jobs=-1,

# Fit the Random Search on the training data
random_search_rf.fit(X_train_pca_df, y_train)

# Print the best parameters found by Random Search
print("Best Parameters (Random Forest):", random_search_rf.best_params_)

# Make predictions on the testing data using the best model
y_pred_random_search_rf = random_search_rf.best_estimator_.predict(X_test_pca_df)

# Evaluate the model after Random Search
accuracy_random_search_rf = accuracy_score(y_test, y_pred_random_search_rf)
precision_random_search_rf = precision_score(y_test, y_pred_random_search_rf, average='weighted')
recall_random_search_rf = recall_score(y_test, y_pred_random_search_rf, average='weighted')
f1_random_search_rf = f1_score(y_test, y_pred_random_search_rf, average='weighted')
conf_matrix_random_search_rf = confusion_matrix(y_test, y_pred_random_search_rf)

# Print the evaluation metrics after Random Search
print("Metrics after Random Search (Random Forest):")
print(f"Accuracy: {accuracy_random_search_rf:.4f}")
print(f"Precision: {precision_random_search_rf:.4f}")
print(f"Recall: {recall_random_search_rf:.4f}")
print(f"F1 Score: {f1_random_search_rf:.4f}")
print("Confusion Matrix:")
print(conf_matrix_random_search_rf)

Best Parameters (Random Forest): {'n_estimators': 100, 'min_weight_fraction_leaf': 0.0, 'min_samples_split': 5, 'min_samples_leaf':
Metrics after Random Search (Random Forest):
Accuracy: 0.6175
Precision: 0.6217
Recall: 0.6175
F1 Score: 0.5846
Confusion Matrix:
[[3349   0  365   68]
 [  18 3118  391  100]
 [ 665 1015 1743  256]
 [ 864 1137  745  870]]

```

Model Evaluation:

Logistic Regression:

Strengths:

Interpretability: Logistic Regression provides interpretable coefficients, allowing you to understand the impact of each feature on the credit mix classification.

Efficiency: It is computationally efficient and may perform well if the relationship between features and credit mix is predominantly linear.

Limitations: *Linear Decision Boundary:* Assumes a linear relationship between features and the target, which might not capture complex patterns in credit mix data.

Sensitivity to Outliers: Logistic Regression can be sensitive to outliers, which may exist in credit-related datasets.

Random Forest:

Strengths:

Non-Linearity: Random Forest can capture non-linear relationships, which may be crucial if the relationship between features and credit mix is complex.

Robustness: Handles outliers and noisy data well, making it suitable for real-world credit-related datasets.

Limitations:

Complexity: The ensemble nature can make it challenging to interpret the model's decision-making process, which might be important in credit analysis.

Potential Overfitting: While Random Forests are less prone to overfitting, they can still be sensitive to noise in the data.

Support Vector Machine (SVM):

Strengths:

Non-Linearity: SVMs can model non-linear relationships, providing flexibility in capturing intricate patterns in credit mix data.

Limitations:

Computational Intensity: SVMs can be computationally intensive, which may be a concern for large credit-related datasets.

Sensitivity to Parameters: The performance of SVMs depends on parameter tuning, and finding the right parameters might require experimentation.

XGBoost:

Strengths:

High Performance: XGBoost often performs well on various tasks, including classification, and might capture complex relationships in credit mix data.

Handling Missing Values: XGBoost can handle missing values internally, which may be beneficial if your dataset has missing information.

Limitations:

Complexity: The complexity of XGBoost models can make them challenging to interpret, potentially limiting their explainability.

Parameter Tuning: Effective use of XGBoost may require careful tuning of hyperparameters, which can be an iterative process.

Model Evaluation Considerations:

Imbalanced Data: If the credit mix classes are imbalanced, consider using evaluation metrics like precision, recall, and F1 score in addition to accuracy.

Interpretability: Consider the interpretability of the models, especially if stakeholders require transparency in credit-related decisions.

Domain Knowledge: Leverage domain knowledge to interpret model results and guide feature engineering.

Interpretability

Logistic Regression:

— .